

Data Preprocessing and Feature Engineering

Module 5, Lab 3: Preparing Data for Machine Learning

Raw data is rarely ready for machine learning algorithms. This lab teaches you how to clean, transform, and engineer features to create high-quality datasets that lead to better model performance.

Learning Objectives

By the end of this lab, you will be able to:

- Handle missing values using various strategies
- Encode categorical variables for machine learning
- Scale and normalize numerical features
- Detect and handle outliers appropriately
- Create new features through feature engineering
- Build preprocessing pipelines for reproducibility

Why This Matters

Data preprocessing often takes 80% of a data scientist's time, but it's crucial for model success. Poor data preparation leads to poor models, regardless of the algorithm used.

Setup and Data Loading

```
# Install required packages
!pip install --upgrade pip
!pip install pandas numpy matplotlib seaborn scikit-learn

Requirement already satisfied: pip in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (25.2)
Requirement already satisfied: pandas in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (2.3.1)
Requirement already satisfied: numpy in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (1.26.4)
Requirement already satisfied: matplotlib in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (3.10.5)
Requirement already satisfied: seaborn in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (0.13.2)
Requirement already satisfied: scikit-learn in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (1.7.1)
Requirement already satisfied: python-dateutil<=2.8.2 in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (from
pandas) (2.9.0.post0)
Requirement already satisfied: pytz<=2020.1 in
```

```
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (from
pandas) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (from
pandas) (2025.2)
Requirement already satisfied: contourpy>=1.0.1 in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (from
matplotlib) (1.3.1)
Requirement already satisfied: cycler>=0.10 in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (from
matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (from
matplotlib) (4.55.0)
Requirement already satisfied: kiwisolver>=1.3.1 in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (from
matplotlib) (1.4.7)
Requirement already satisfied: packaging>=20.0 in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (from
matplotlib) (25.0)
Requirement already satisfied: pillow>=8 in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (from
matplotlib) (11.0.0)
Requirement already satisfied: pyparsing>=2.3.1 in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (from
matplotlib) (3.0.9)
Requirement already satisfied: scipy>=1.8.0 in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (from
scikit-learn) (1.13.1)
Requirement already satisfied: joblib>=1.2.0 in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (from
scikit-learn) (1.5.1)
Requirement already satisfied: threadpoolctl>=3.1.0 in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (from
scikit-learn) (3.5.0)
Requirement already satisfied: six>=1.5 in
/Users/martin.demel/myenv3.10/lib/python3.10/site-packages (from
python-dateutil>=2.8.2->pandas) (1.16.0)
```

```
# Import necessary libraries
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler, MinMaxScaler,
LabelEncoder, OneHotEncoder
from sklearn.impute import SimpleImputer, KNNImputer
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.pipeline import Pipeline
```

```

from sklearn.compose import ColumnTransformer
import warnings
warnings.filterwarnings('ignore')

# Set plotting style
plt.style.use('seaborn-v0_8')
%matplotlib inline

print("Libraries imported successfully!")

Libraries imported successfully!

```

Creating a Realistic Dataset with Data Quality Issues

We'll create a dataset that mimics real-world data problems you'll encounter.

```

# Create a realistic employee dataset with various data quality issues
np.random.seed(42)
n_employees = 1000

# Generate base employee data
employee_data = {
    'employee_id': range(1, n_employees + 1),
    'age': np.random.normal(35, 10, n_employees),
    'years_experience': np.random.exponential(5, n_employees),
    'education_level': np.random.choice(['High School', 'Bachelor',
'Master', 'PhD'],
                                         n_employees, p=[0.2, 0.5, 0.25,
0.05]),
    'department': np.random.choice(['Engineering', 'Sales',
'Marketing', 'HR', 'Finance'],
                                     n_employees, p=[0.3, 0.25, 0.2,
0.15, 0.1]),
    'job_level': np.random.choice(['Junior', 'Mid', 'Senior', 'Lead'],
                                   n_employees, p=[0.3, 0.4, 0.25,
0.05]),
    'location': np.random.choice(['New York', 'San Francisco',
'Chicago', 'Austin', 'Remote'],
                                  n_employees, p=[0.25, 0.2, 0.15, 0.15,
0.25]),
    'performance_score': np.random.normal(7.5, 1.5, n_employees),
    'hours_per_week': np.random.normal(42, 8, n_employees),
    'projects_completed': np.random.poisson(8, n_employees),
    'training_hours': np.random.gamma(2, 10, n_employees)
}

# Create DataFrame
df = pd.DataFrame(employee_data)

```

```

# Add realistic constraints
df['age'] = np.clip(df['age'], 22, 65)
df['years_experience'] = np.clip(df['years_experience'], 0, df['age']
- 22)
df['performance_score'] = np.clip(df['performance_score'], 1, 10)
df['hours_per_week'] = np.clip(df['hours_per_week'], 20, 60)

# Create salary based on realistic factors (this will be our target
variable)
base_salary = 50000
education_bonus = {'High School': 0, 'Bachelor': 15000, 'Master':
25000, 'PhD': 40000}
level_bonus = {'Junior': 0, 'Mid': 20000, 'Senior': 40000, 'Lead':
70000}
dept_bonus = {'Engineering': 15000, 'Sales': 10000, 'Marketing': 5000,
'HR': 0, 'Finance': 8000}
location_bonus = {'New York': 20000, 'San Francisco': 25000,
'Chicago': 5000, 'Austin': 8000, 'Remote': 0}

df['salary'] = (base_salary +
df['education_level'].map(education_bonus) +
df['job_level'].map(level_bonus) +
df['department'].map(dept_bonus) +
df['location'].map(location_bonus) +
df['years_experience'] * 2000 +
df['performance_score'] * 3000 +
np.random.normal(0, 10000, n_employees))

df['salary'] = np.maximum(df['salary'], 35000) # Minimum salary

print(f"Dataset created with {len(df)} employees")
print(f"Dataset shape: {df.shape}")
df.head()

```

Dataset created with 1000 employees
Dataset shape: (1000, 12)

	employee_id	age	years_experience	education_level
0	1	39.967142	0.916506	Bachelor
Marketing				
1	2	33.617357	0.552244	High School
Engineering				
2	3	41.476885	5.058921	High School
Sales				
3	4	50.230299	6.128975	Bachelor
HR				
4	5	32.658466	0.160479	Master
Sales				

	job_level	location	performance_score	hours_per_week	\
0	Junior	Chicago	6.516468	52.944635	
1	Junior	Austin	7.146433	43.007211	
2	Mid	San Francisco	6.090805	48.808993	
3	Mid	New York	9.178867	51.790847	
4	Junior	Remote	7.058316	44.703644	

	projects_completed	training_hours	salary
0	5	15.103473	98155.118139
1	6	10.621985	98464.144788
2	7	81.025460	160878.080717
3	8	26.356775	145460.439793
4	5	26.094782	107113.789261

Introducing Realistic Data Quality Issues

```
# Introduce missing values (realistic patterns)
# Performance scores might be missing for new employees
new_employee_mask = df['years_experience'] < 0.5
df.loc[new_employee_mask & (np.random.random(len(df)) < 0.3),
'performance_score'] = np.nan

# Training hours might be missing randomly
missing_training = np.random.choice(df.index, size=80, replace=False)
df.loc[missing_training, 'training_hours'] = np.nan

# Some education levels might be missing
missing_education = np.random.choice(df.index, size=30, replace=False)
df.loc[missing_education, 'education_level'] = np.nan

# Add some outliers
# Extremely high performers
outlier_indices = np.random.choice(df.index, size=10, replace=False)
df.loc[outlier_indices, 'hours_per_week'] = np.random.uniform(70, 80,
10)
df.loc[outlier_indices, 'projects_completed'] = np.random.uniform(25,
35, 10)

# Add some inconsistent data
# Some employees with PhD but very low experience (career changers)
career_changer_indices = np.random.choice(df[df['education_level'] ==
'PhD'].index, size=5, replace=False)
df.loc[career_changer_indices, 'years_experience'] =
np.random.uniform(0, 2, 5)

# Add some duplicate-like entries (same person, different records)
duplicate_base = df.sample(3).copy()
duplicate_base['employee_id'] = range(n_employees + 1, n_employees +
4)
# Slightly modify some values to simulate data entry errors
```

```

duplicate_base['age'] += np.random.randint(-1, 2, 3)
duplicate_base['salary'] += np.random.randint(-5000, 5000, 3)
df = pd.concat([df, duplicate_base], ignore_index=True)

print("Data quality issues introduced:")
print(f"Missing values: {df.isnull().sum().sum()}")
print(f"Total records: {len(df)}")
print("\nMissing values by column:")
print(df.isnull().sum()[df.isnull().sum() > 0])

Data quality issues introduced:
Missing values: 150
Total records: 1003

Missing values by column:
education_level      30
performance_score    40
training_hours       80
dtype: int64

```

Step 1: Data Quality Assessment

Before preprocessing, let's understand what we're working with.

```

# Comprehensive data quality report
def data_quality_report(df):
    print("=== DATA QUALITY REPORT ===")
    print(f"\n Dataset Overview:")
    print(f"    • Shape: {df.shape}")
    print(f"    • Memory usage: {df.memory_usage(deep=True).sum() / 1024:.1f} KB")

    print(f"\n Missing Values:")
    missing_data = df.isnull().sum()
    missing_percent = (missing_data / len(df)) * 100

    for col in missing_data[missing_data > 0].index:
        print(f"    • {col}: {missing_data[col]} ({missing_percent[col]:.1f}%)")

    print(f"\n Data Types:")
    print(f"    • Numerical columns: {len(df.select_dtypes(include=[np.number]).columns)}")
    print(f"    • Categorical columns: {len(df.select_dtypes(include=['object']).columns)}")

    print(f"\n Potential Issues:")
    # Check for duplicates
    duplicates = len(df) - len(df.drop_duplicates())

```

```

if duplicates > 0:
    print(f"    • Duplicate rows: {duplicates}")

# Check for outliers in numerical columns
numerical_cols = df.select_dtypes(include=[np.number]).columns
for col in numerical_cols:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1
    outliers = len(df[(df[col] < Q1 - 1.5*IQR) | (df[col] > Q3 +
1.5*IQR)])
    if outliers > 0:
        print(f"    • {col} outliers: {outliers}
({outliers/len(df)*100:.1f}%)")

data_quality_report(df)

=== DATA QUALITY REPORT ===

□ Dataset Overview:
    • Shape: (1003, 12)
    • Memory usage: 312.6 KB

□ Missing Values:
    • education_level: 30 (3.0%)
    • performance_score: 40 (4.0%)
    • training_hours: 80 (8.0%)

□ Data Types:
    • Numerical columns: 8
    • Categorical columns: 4

□ Potential Issues:
    • age outliers: 4 (0.4%)
    • years_experience outliers: 34 (3.4%)
    • performance_score outliers: 3 (0.3%)
    • hours_per_week outliers: 10 (1.0%)
    • projects_completed outliers: 11 (1.1%)
    • training_hours outliers: 10 (1.0%)
    • salary outliers: 6 (0.6%)

# Visualize missing data patterns
plt.figure(figsize=(12, 8))

# Missing data heatmap
plt.subplot(2, 2, 1)
sns.heatmap(df.isnull(), cbar=True, yticklabels=False, cmap='viridis')
plt.title('Missing Data Pattern')

# Missing data bar plot

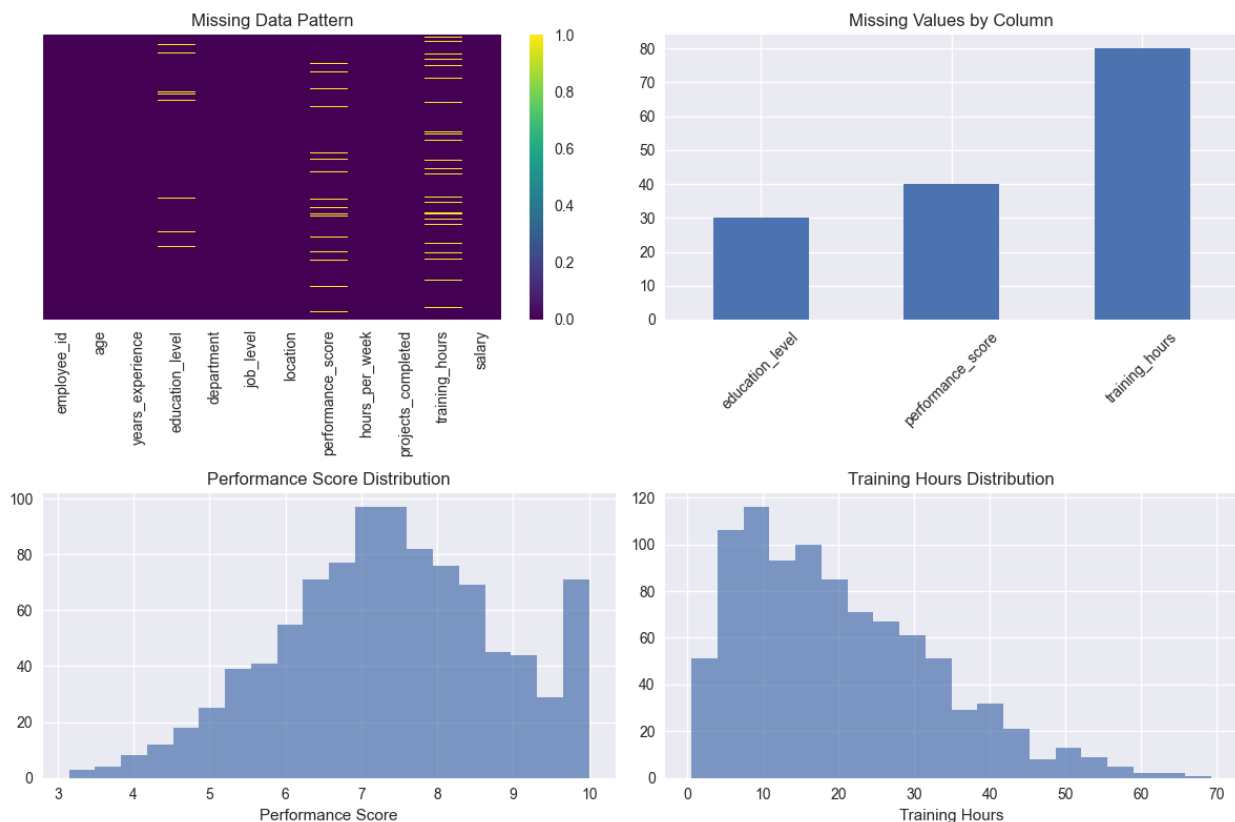
```

```
plt.subplot(2, 2, 2)
missing_counts = df.isnull().sum()
missing_counts = missing_counts[missing_counts > 0]
missing_counts.plot(kind='bar')
plt.title('Missing Values by Column')
plt.xticks(rotation=45)

# Distribution of numerical variables with missing values
plt.subplot(2, 2, 3)
df['performance_score'].hist(bins=20, alpha=0.7)
plt.title('Performance Score Distribution')
plt.xlabel('Performance Score')

plt.subplot(2, 2, 4)
df['training_hours'].hist(bins=20, alpha=0.7)
plt.title('Training Hours Distribution')
plt.xlabel('Training Hours')

plt.tight_layout()
plt.show()
```



Step 2: Handling Missing Values

Different strategies work better for different types of missing data.

2.1 Understanding Missing Data Patterns

```
# Analyze missing data patterns
print("Missing Data Analysis:")
print("\n1. Performance Score Missing Pattern:")
missing_perf = df[df['performance_score'].isnull()]
print(f"    • Average years of experience:
{missing_perf['years_experience'].mean():.2f}")
print(f"    • Most common job level:
{missing_perf['job_level'].mode().iloc[0]}")

print("\n2. Training Hours Missing Pattern:")
missing_training = df[df['training_hours'].isnull()]
print(f"    • Average age: {missing_training['age'].mean():.2f}")
print(f"    • Department distribution:")
print(missing_training['department'].value_counts())

print("\n3. Education Level Missing Pattern:")
missing_edu = df[df['education_level'].isnull()]
print(f"    • Average salary: ${missing_edu['salary'].mean():.0f}")
print(f"    • Average years experience:
{missing_edu['years_experience'].mean():.2f}")
```

Missing Data Analysis:

1. Performance Score Missing Pattern:

- Average years of experience: 0.17
- Most common job level: Junior

2. Training Hours Missing Pattern:

- Average age: 35.49
- Department distribution:

department

Engineering 21

Sales 20

Marketing 15

HR 13

Finance 11

Name: count, dtype: int64

3. Education Level Missing Pattern:

- Average salary: \$137495
- Average years experience: 3.90

2.2 Imputation Strategies

```
# Create a copy for preprocessing
df_processed = df.copy()

print("Applying different imputation strategies...")
```

```

# Strategy 1: Mean imputation for performance_score (numerical)
# But let's be smarter - use group mean based on job level
performance_means = df_processed.groupby('job_level')
['performance_score'].mean()
print("\nPerformance score means by job level:")
print(performance_means)

for level in performance_means.index:
    mask = (df_processed['job_level'] == level) &
(df_processed['performance_score'].isnull())
    df_processed.loc[mask, 'performance_score'] =
performance_means[level]

# Strategy 2: Median imputation for training_hours (skewed
distribution)
training_median = df_processed['training_hours'].median()
df_processed['training_hours'].fillna(training_median, inplace=True)
print(f"\nFilled training hours with median: {training_median:.1f}")

# Strategy 3: Mode imputation for education_level (categorical)
education_mode = df_processed['education_level'].mode().iloc[0]
df_processed['education_level'].fillna(education_mode, inplace=True)
print(f"Filled education level with mode: {education_mode}")

print(f"\nMissing values after imputation:
{df_processed.isnull().sum().sum()}")

```

Applying different imputation strategies...

```

Performance score means by job level:
job_level
Junior      7.483550
Lead        7.326790
Mid         7.256741
Senior      7.435315
Name: performance_score, dtype: float64

```

```

Filled training hours with median: 17.6
Filled education level with mode: Bachelor

```

Missing values after imputation: 0

```

# Compare distributions before and after imputation
fig, axes = plt.subplots(2, 2, figsize=(15, 10))

```

```

# Performance score
axes[0, 0].hist(df['performance_score'].dropna(), bins=20, alpha=0.7,
label='Original', density=True)
axes[0, 0].hist(df_processed['performance_score'], bins=20, alpha=0.7,
label='After Imputation', density=True)

```

```

axes[0, 0].set_title('Performance Score Distribution')
axes[0, 0].legend()

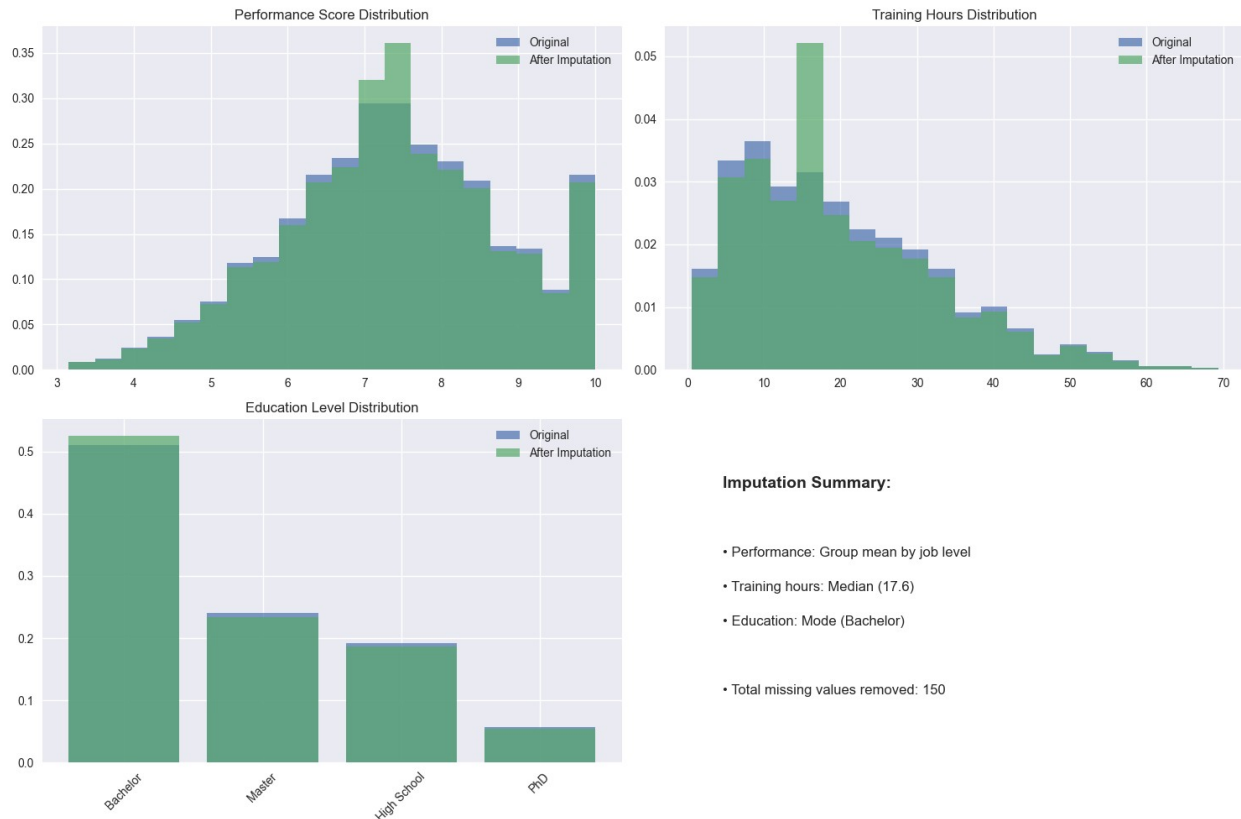
# Training hours
axes[0, 1].hist(df['training_hours'].dropna(), bins=20, alpha=0.7,
label='Original', density=True)
axes[0, 1].hist(df_processed['training_hours'], bins=20, alpha=0.7,
label='After Imputation', density=True)
axes[0, 1].set_title('Training Hours Distribution')
axes[0, 1].legend()

# Education level
original_edu = df['education_level'].value_counts(normalize=True)
imputed_edu =
df_processed['education_level'].value_counts(normalize=True)
axes[1, 0].bar(range(len(original_edu)), original_edu.values,
alpha=0.7, label='Original')
axes[1, 0].bar(range(len(imputed_edu)), imputed_edu.values, alpha=0.7,
label='After Imputation')
axes[1, 0].set_xticks(range(len(original_edu)))
axes[1, 0].set_xticklabels(original_edu.index, rotation=45)
axes[1, 0].set_title('Education Level Distribution')
axes[1, 0].legend()

# Summary
axes[1, 1].text(0.1, 0.8, 'Imputation Summary:', fontsize=14,
fontweight='bold')
axes[1, 1].text(0.1, 0.6, f'• Performance: Group mean by job level',
fontsize=12)
axes[1, 1].text(0.1, 0.5, f'• Training hours: Median
({training_median:.1f})', fontsize=12)
axes[1, 1].text(0.1, 0.4, f'• Education: Mode ({education_mode})',
fontsize=12)
axes[1, 1].text(0.1, 0.2, f'• Total missing values removed:
{df.isnull().sum().sum()}', fontsize=12)
axes[1, 1].set_xlim(0, 1)
axes[1, 1].set_ylim(0, 1)
axes[1, 1].axis('off')

plt.tight_layout()
plt.show()

```



Step 3: Encoding Categorical Variables

Machine learning algorithms work with numbers, so we need to convert categorical data.

```
# Identify categorical columns
categorical_columns = ['education_level', 'department', 'job_level',
                       'location']
print("Categorical columns to encode:")
for col in categorical_columns:
    print(f"• {col}: {df_processed[col].nunique()} unique values")
    print(f"  Values: {list(df_processed[col].unique())}")
    print()
```

Categorical columns to encode:

- education_level: 4 unique values
Values: ['Bachelor', 'High School', 'Master', 'PhD']
- department: 5 unique values
Values: ['Marketing', 'Engineering', 'Sales', 'HR', 'Finance']
- job_level: 4 unique values
Values: ['Junior', 'Mid', 'Senior', 'Lead']
- location: 5 unique values

```
Values: ['Chicago', 'Austin', 'San Francisco', 'New York', 'Remote']
```

3.1 Ordinal Encoding (for ordered categories)

```
# Education level has a natural order
education_order = ['High School', 'Bachelor', 'Master', 'PhD']
education_mapping = {level: i for i, level in
                     enumerate(education_order)}
df_processed['education_level_encoded'] =
df_processed['education_level'].map(education_mapping)

# Job level also has a natural order
job_order = ['Junior', 'Mid', 'Senior', 'Lead']
job_mapping = {level: i for i, level in enumerate(job_order)}
df_processed['job_level_encoded'] =
df_processed['job_level'].map(job_mapping)

print("Ordinal Encoding Applied:")
print("\nEducation Level Mapping:")
for original, encoded in education_mapping.items():
    print(f" {original} → {encoded}")

print("\nJob Level Mapping:")
for original, encoded in job_mapping.items():
    print(f" {original} → {encoded}")
```

Ordinal Encoding Applied:

Education Level Mapping:

High School → 0
Bachelor → 1
Master → 2
PhD → 3

Job Level Mapping:

Junior → 0
Mid → 1
Senior → 2
Lead → 3

3.2 One-Hot Encoding (for nominal categories)

```
# Department and location don't have natural order - use one-hot
encoding
# Create dummy variables
department_dummies = pd.get_dummies(df_processed['department'],
prefix='dept')
location_dummies = pd.get_dummies(df_processed['location'],
prefix='loc')
```

```
print("One-Hot Encoding Applied:")
print(f"\nDepartment columns created:
{list(department_dummies.columns)}")
print(f"Location columns created: {list(location_dummies.columns)}")
```

```
# Add to dataframe
```

```
df_processed = pd.concat([df_processed, department_dummies,
location_dummies], axis=1)
```

```
print(f"\nDataset shape after encoding: {df_processed.shape}")
```

One-Hot Encoding Applied:

Department columns created: ['dept_Engineering', 'dept_Finance',
'dept_HR', 'dept_Marketing', 'dept_Sales']
Location columns created: ['loc_Austin', 'loc_Chicago', 'loc_New
York', 'loc_Remote', 'loc_San Francisco']

Dataset shape after encoding: (1003, 24)

```
# Visualize the encoding results
```

```
fig, axes = plt.subplots(2, 2, figsize=(15, 10))
```

```
# Education level encoding
```

```
education_comparison = pd.DataFrame({
    'Original': df_processed['education_level'].value_counts(),
    'Encoded': df_processed.groupby('education_level')
['education_level_encoded'].first()
})
```

```
axes[0, 0].bar(education_comparison.index,
education_comparison['Encoded'])
axes[0, 0].set_title('Education Level Ordinal Encoding')
axes[0, 0].set_ylabel('Encoded Value')
axes[0, 0].tick_params(axis='x', rotation=45)
```

```
# Job level encoding
```

```
job_comparison = pd.DataFrame({
    'Original': df_processed['job_level'].value_counts(),
    'Encoded': df_processed.groupby('job_level')
['job_level_encoded'].first()
})
```

```
axes[0, 1].bar(job_comparison.index, job_comparison['Encoded'])
axes[0, 1].set_title('Job Level Ordinal Encoding')
axes[0, 1].set_ylabel('Encoded Value')
axes[0, 1].tick_params(axis='x', rotation=45)
```

```
# Department one-hot encoding
```

```
dept_cols = [col for col in df_processed.columns if
col.startswith('dept_')]
```

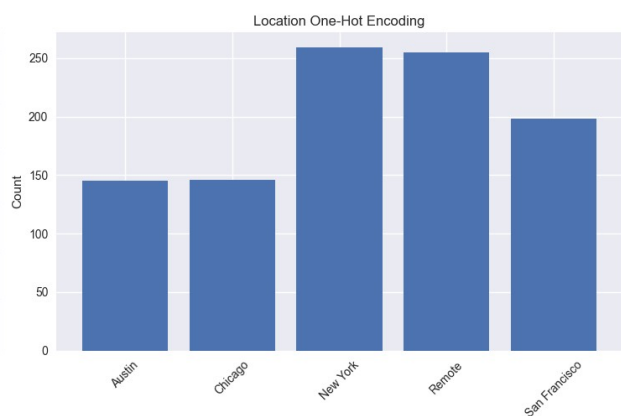
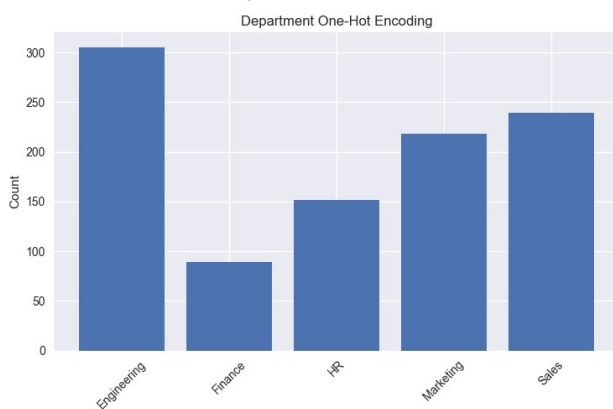
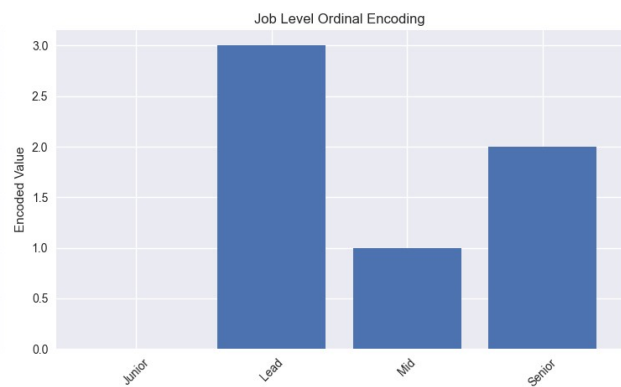
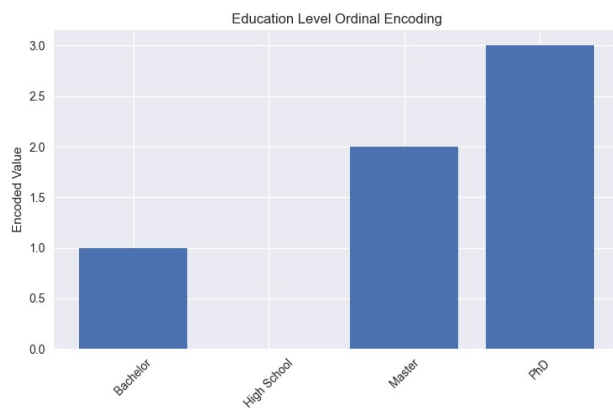
```

dept_sums = df_processed[dept_cols].sum()
axes[1, 0].bar(range(len(dept_sums)), dept_sums.values)
axes[1, 0].set_xticks(range(len(dept_sums)))
axes[1, 0].set_xticklabels([col.replace('dept_', '') for col in
dept_sums.index], rotation=45)
axes[1, 0].set_title('Department One-Hot Encoding')
axes[1, 0].set_ylabel('Count')

# Location one-hot encoding
loc_cols = [col for col in df_processed.columns if
col.startswith('loc_')]
loc_sums = df_processed[loc_cols].sum()
axes[1, 1].bar(range(len(loc_sums)), loc_sums.values)
axes[1, 1].set_xticks(range(len(loc_sums)))
axes[1, 1].set_xticklabels([col.replace('loc_', '') for col in
loc_sums.index], rotation=45)
axes[1, 1].set_title('Location One-Hot Encoding')
axes[1, 1].set_ylabel('Count')

plt.tight_layout()
plt.show()

```



Step 4: Feature Scaling and Normalization

Different features have different scales, which can bias machine learning algorithms.

```
# Identify numerical columns for scaling
numerical_columns = ['age', 'years_experience', 'performance_score',
                     'hours_per_week',
                     'projects_completed', 'training_hours', 'salary']

print("Numerical columns statistics before scaling:")
print(df_processed[numerical_columns].describe().round(2))
```

Numerical columns statistics before scaling:

	age	years_experience	performance_score	hours_per_week	\
count	1003.00	1003.00	1003.00	1003.00	
mean	35.52	3.90	7.37	42.09	
std	9.11	3.79	1.42	8.73	
min	22.00	0.00	3.15	20.30	
25%	28.49	0.86	6.47	36.00	
50%	35.25	2.93	7.42	41.74	
75%	41.47	5.93	8.33	47.69	
max	65.00	19.34	10.00	78.85	

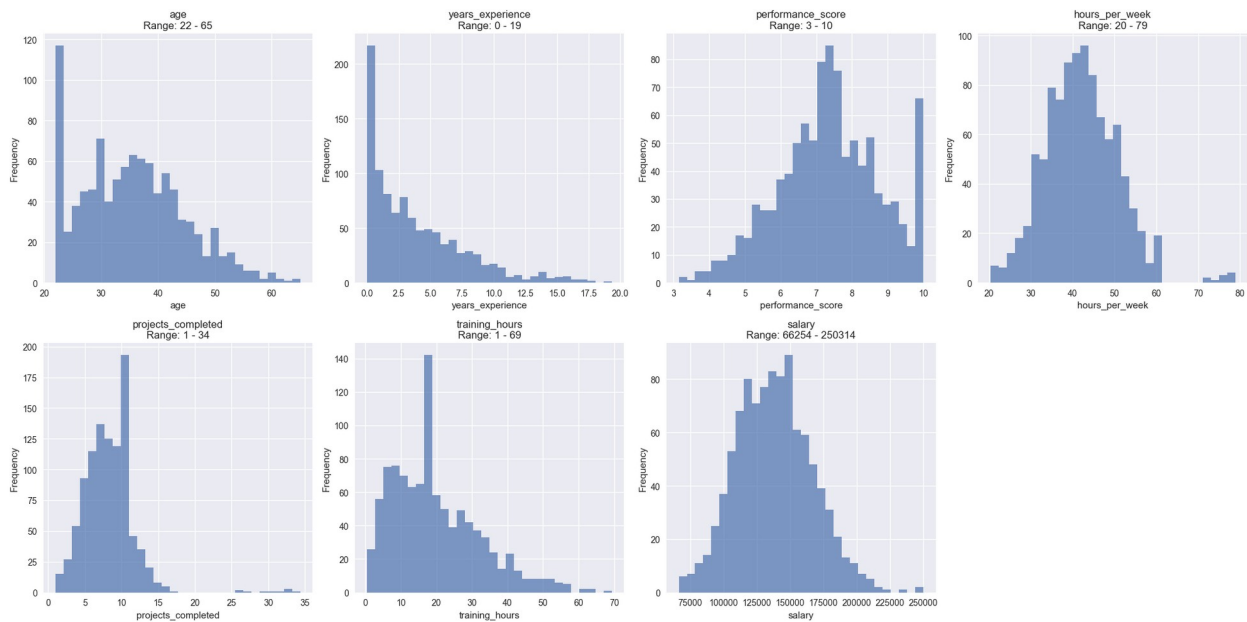
	projects_completed	training_hours	salary
count	1003.00	1003.00	1003.00
mean	8.30	19.86	138370.24
std	3.60	12.40	28602.46
min	1.00	0.52	66253.73
25%	6.00	10.23	117935.23
50%	8.00	17.61	137586.00
75%	10.00	27.27	157437.15
max	34.36	69.34	250314.44

```
# Visualize the scale differences
fig, axes = plt.subplots(2, 4, figsize=(20, 10))
axes = axes.ravel()

for i, col in enumerate(numerical_columns):
    axes[i].hist(df_processed[col], bins=30, alpha=0.7)
    axes[i].set_title(f'{col}\nRange: {df_processed[col].min():.0f} - {df_processed[col].max():.0f}')
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('Frequency')

# Remove empty subplot
fig.delaxes(axes[7])

plt.tight_layout()
plt.show()
```

4.1 Standard Scaling (Z-score normalization)

```
# Apply StandardScaler (mean=0, std=1)
scaler_standard = StandardScaler()

# We'll exclude salary from scaling since it's our target variable
features_to_scale = [col for col in numerical_columns if col != 'salary']

# Fit and transform
df_standard_scaled = df_processed.copy()
df_standard_scaled[features_to_scale] =
scaler_standard.fit_transform(df_processed[features_to_scale])

print("Standard Scaling Applied:")
print("\nFeatures scaled (mean=0, std=1):")
print(df_standard_scaled[features_to_scale].describe().round(3))
```

Standard Scaling Applied:

Features scaled (mean=0, std=1):

	age	years_experience	performance_score	hours_per_week
count	1003.000	1003.000	1003.000	1003.000
mean	-0.000	-0.000	0.000	0.000
std	1.000	1.000	1.000	1.000
min	-1.484	-1.031	-2.976	-2.497
25%	-0.772	-0.803	-0.636	-0.698

50%	-0.030	-0.257	0.032	-0.040
75%	0.653	0.536	0.673	0.641
max	3.237	4.080	1.852	4.211

	projects_completed	training_hours
count	1003.000	1003.000
mean	0.000	0.000
std	1.000	1.000
min	-2.032	-1.561
25%	-0.640	-0.777
50%	-0.084	-0.182
75%	0.473	0.597
max	7.252	3.994

4.2 Min-Max Scaling (0-1 normalization)

```
# Apply MinMaxScaler (range 0-1)
```

```
scaler_minmax = MinMaxScaler()
```

```
# Fit and transform
```

```
df_minmax_scaled = df_processed.copy()
```

```
df_minmax_scaled[features_to_scale] =
```

```
scaler_minmax.fit_transform(df_processed[features_to_scale])
```

```
print("Min-Max Scaling Applied:")
```

```
print("\nFeatures scaled (range 0-1):")
```

```
print(df_minmax_scaled[features_to_scale].describe().round(3))
```

Min-Max Scaling Applied:

Features scaled (range 0-1):

	age	years_experience	performance_score	hours_per_week
count	1003.000	1003.000	1003.000	1003.000
mean	0.314	0.202	0.616	0.372
std	0.212	0.196	0.207	0.149
min	0.000	0.000	0.000	0.000
25%	0.151	0.045	0.485	0.268
50%	0.308	0.151	0.623	0.366
75%	0.453	0.307	0.756	0.468

max	1.000	1.000	1.000	1.000
-----	-------	-------	-------	-------

	projects_completed	training_hours
count	1003.000	1003.000
mean	0.219	0.281
std	0.108	0.180
min	0.000	0.000
25%	0.150	0.141
50%	0.210	0.248
75%	0.270	0.389
max	1.000	1.000

```
# Compare scaling methods
```

```
fig, axes = plt.subplots(3, 3, figsize=(18, 15))
```

```
# Select a few key features for comparison
```

```
comparison_features = ['age', 'years_experience', 'salary']
```

```
for i, feature in enumerate(comparison_features):
```

```
    # Original
```

```
    axes[0, i].hist(df_processed[feature], bins=30, alpha=0.7,
color='blue')
```

```
    axes[0, i].set_title(f'Original {feature}')
```

```
    axes[0, i].set_ylabel('Frequency')
```

```
    if feature != 'salary': # Don't scale target variable
```

```
        # Standard scaled
```

```
        axes[1, i].hist(df_standard_scaled[feature], bins=30,
alpha=0.7, color='green')
```

```
        axes[1, i].set_title(f'Standard Scaled {feature}')
```

```
        axes[1, i].set_ylabel('Frequency')
```

```
        # Min-max scaled
```

```
        axes[2, i].hist(df_minmax_scaled[feature], bins=30, alpha=0.7,
color='red')
```

```
        axes[2, i].set_title(f'Min-Max Scaled {feature}')
```

```
        axes[2, i].set_ylabel('Frequency')
```

```
    else:
```

```
        # For salary, show the same distribution
```

```
        axes[1, i].hist(df_processed[feature], bins=30, alpha=0.7,
color='blue')
```

```
        axes[1, i].set_title(f'{feature} (Target - Not Scaled)')
```


```
        axes[1, i].set_ylabel('Frequency')
```

```
        axes[2, i].hist(df_processed[feature], bins=30, alpha=0.7,
color='blue')
```

```
        axes[2, i].set_title(f'{feature} (Target - Not Scaled)')
```

```
        axes[2, i].set_ylabel('Frequency')
```

```
plt.tight_layout()
plt.show()
```



Step 5: Outlier Detection and Treatment

Outliers can significantly impact model performance.

```
# Detect outliers using IQR method
def detect_outliers_iqr(df, column, multiplier=1.5):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - multiplier * IQR
    upper_bound = Q3 + multiplier * IQR

    outliers = df[(df[column] < lower_bound) | (df[column] >
upper_bound)]
```

```

    return outliers, lower_bound, upper_bound

# Detect outliers in key numerical columns
outlier_columns = ['hours_per_week', 'projects_completed', 'salary']

print("Outlier Detection Results:")
for col in outlier_columns:
    outliers, lower, upper = detect_outliers_iqr(df_processed, col)
    print(f"\n{col}:")
    print(f"    Normal range: {lower:.2f} to {upper:.2f}")
    print(f"    Outliers found: {len(outliers)}")
    print(f"    ({len(outliers)/len(df_processed)*100:.1f}%)")
    if len(outliers) > 0:
        print(f"    Outlier values: {sorted(outliers[col].values)
[:5]}...) # Show first 5

Outlier Detection Results:

hours_per_week:
    Normal range: 18.46 to 65.22
    Outliers found: 10 (1.0%)
    Outlier values: [71.94835522592277, 72.7064801321802,
73.86425859846484, 75.22909034668608, 76.49865962409]...

projects_completed:
    Normal range: 0.00 to 16.00
    Outliers found: 11 (1.1%)
    Outlier values: [17.0, 26.124664359613647, 26.308214262651816,
26.589054348754964, 29.42997696754455]...

salary:
    Normal range: 58682.34 to 216690.04
    Outliers found: 6 (0.6%)
    Outlier values: [218159.66297778895, 218796.9428226131,
221022.70465305357, 235655.16665500513, 249670.1677638195]...

# Visualize outliers
fig, axes = plt.subplots(2, 3, figsize=(18, 12))
axes = axes.ravel()

for i, col in enumerate(outlier_columns):
    # Box plot
    axes[i*2].boxplot(df_processed[col])
    axes[i*2].set_title(f'{col} - Box Plot')
    axes[i*2].set_ylabel(col)

    # Histogram with outlier boundaries
    outliers, lower, upper = detect_outliers_iqr(df_processed, col)
    axes[i*2+1].hist(df_processed[col], bins=30, alpha=0.7)
    axes[i*2+1].axvline(lower, color='red', linestyle='--',

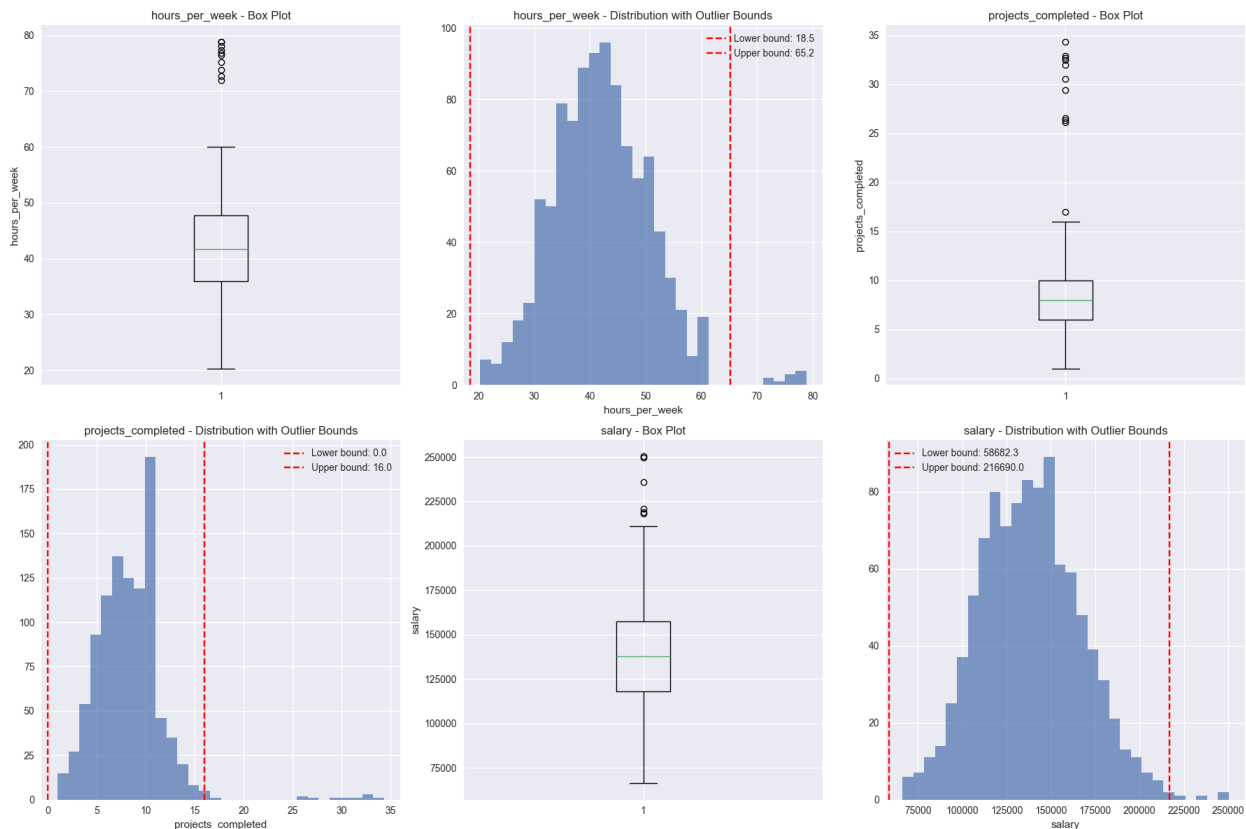
```

```

label=f'Lower bound: {lower:.1f}'
axes[i*2+1].axvline(upper, color='red', linestyle='--',
label=f'Upper bound: {upper:.1f}')
axes[i*2+1].set_title(f'{col} - Distribution with Outlier Bounds')
axes[i*2+1].set_xlabel(col)
axes[i*2+1].legend()

plt.tight_layout()
plt.show()

```



```

# Handle outliers - we'll use capping (Winsorization)
df_outlier_treated = df_processed.copy()

for col in outlier_columns:
    outliers, lower, upper = detect_outliers_iqr(df_processed, col)

    # Cap outliers at the bounds
    df_outlier_treated[col] =
df_outlier_treated[col].clip(lower=lower, upper=upper)

    print(f"\n{col} outlier treatment:")
    print(f"  Values capped below {lower:.2f}:
{len(df_processed[df_processed[col] < lower])}")
    print(f"  Values capped above {upper:.2f}:

```

```
{len(df_processed[df_processed[col] > upper])}"")
print("\nOutlier treatment completed using Winsorization (capping).")

hours_per_week outlier treatment:
  Values capped below 18.46: 0
  Values capped above 65.22: 10

projects_completed outlier treatment:
  Values capped below 0.00: 0
  Values capped above 16.00: 11

salary outlier treatment:
  Values capped below 58682.34: 0
  Values capped above 216690.04: 6

Outlier treatment completed using Winsorization (capping).
```

Step 6: Feature Engineering

Creating new features that might be more predictive than the original ones.

```
# Create new features based on domain knowledge
df_engineered = df_outlier_treated.copy()

print("Creating new features...")

# 1. Experience-to-age ratio (career focus indicator)
df_engineered['experience_age_ratio'] =
df_engineered['years_experience'] / df_engineered['age']

# 2. Productivity score (projects per year of experience)
df_engineered['productivity_score'] =
df_engineered['projects_completed'] /
(df_engineered['years_experience'] + 1) # +1 to avoid division by
zero

# 3. Work intensity (hours per week relative to standard 40)
df_engineered['work_intensity'] = df_engineered['hours_per_week'] / 40

# 4. Training investment (training hours per year of experience)
df_engineered['training_investment'] = df_engineered['training_hours']
/ (df_engineered['years_experience'] + 1)

# 5. Performance-experience interaction
df_engineered['performance_experience'] =
df_engineered['performance_score'] * df_engineered['years_experience']

# 6. Age groups (categorical feature from numerical)
```

```

df_engineered['age_group'] = pd.cut(df_engineered['age'],
                                     bins=[0, 30, 40, 50, 100],
                                     labels=['Young', 'Mid-Career',
'Experienced', 'Senior'])

# 7. Experience level (categorical feature from numerical)
df_engineered['experience_level'] =
pd.cut(df_engineered['years_experience'],
        bins=[0, 2, 5, 10, 100],
        labels=['Novice',
'Intermediate', 'Experienced', 'Expert'])

# 8. High performer flag
df_engineered['high_performer'] = (df_engineered['performance_score']
> df_engineered['performance_score'].quantile(0.75)).astype(int)

print("New features created:")
new_features = ['experience_age_ratio', 'productivity_score',
'work_intensity',
'training_investment', 'performance_experience',
'age_group',
'experience_level', 'high_performer']
for feature in new_features:
    print(f" • {feature}")

print(f"\nDataset shape after feature engineering:
{df_engineered.shape}")

Creating new features...
New features created:
  • experience_age_ratio
  • productivity_score
  • work_intensity
  • training_investment
  • performance_experience
  • age_group
  • experience_level
  • high_performer

Dataset shape after feature engineering: (1003, 32)

# Analyze the new features
fig, axes = plt.subplots(2, 4, figsize=(20, 10))
axes = axes.ravel()

numerical_new_features = ['experience_age_ratio',
'productivity_score', 'work_intensity',
'training_investment',
'performance_experience']

```



```

for i, feature in enumerate(numerical_new_features):
    axes[i].hist(df_engineered[feature], bins=30, alpha=0.7)
    axes[i].set_title(f'{feature}')
    axes[i].set_xlabel(feature)
    axes[i].set_ylabel('Frequency')

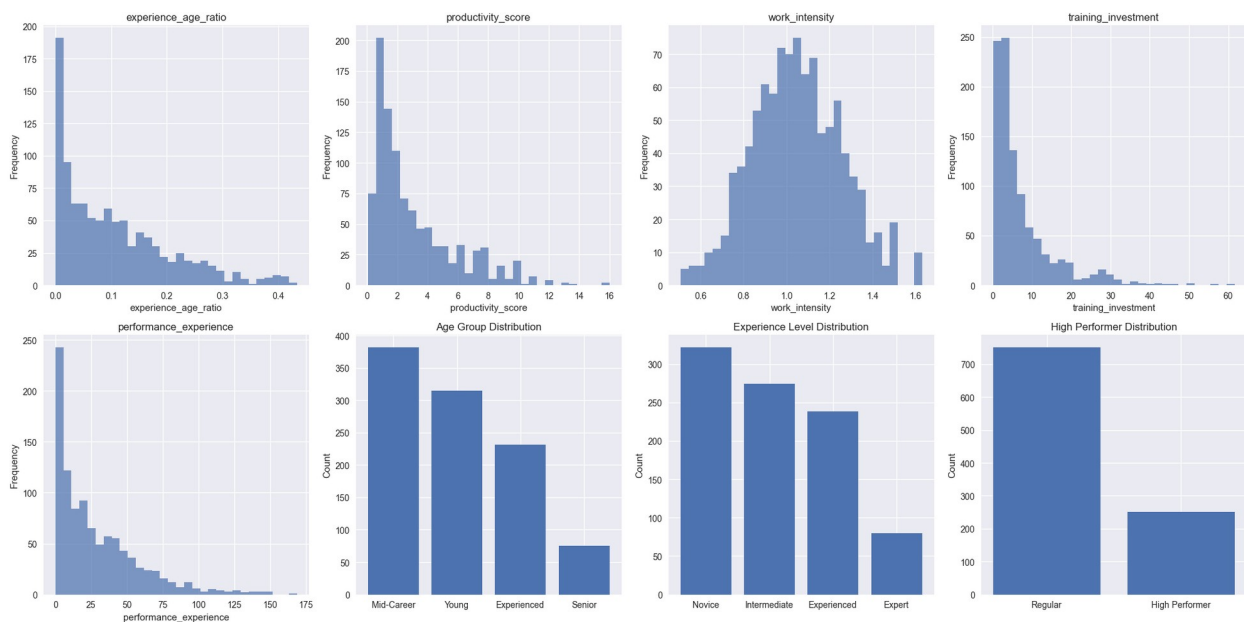
# Age group distribution
age_group_counts = df_engineered['age_group'].value_counts()
axes[5].bar(age_group_counts.index, age_group_counts.values)
axes[5].set_title('Age Group Distribution')
axes[5].set_ylabel('Count')

# Experience level distribution
exp_level_counts = df_engineered['experience_level'].value_counts()
axes[6].bar(exp_level_counts.index, exp_level_counts.values)
axes[6].set_title('Experience Level Distribution')
axes[6].set_ylabel('Count')

# High performer distribution
high_perf_counts = df_engineered['high_performer'].value_counts()
axes[7].bar(['Regular', 'High Performer'], high_perf_counts.values)
axes[7].set_title('High Performer Distribution')
axes[7].set_ylabel('Count')

plt.tight_layout()
plt.show()

```



Step 7: Feature Selection

Not all features are equally important. Let's identify the most predictive ones.

```

# Prepare data for feature importance analysis
# Select numerical features for correlation analysis
numerical_features_all =
df_engineered.select_dtypes(include=[np.number]).columns.tolist()
# Remove ID and target variable
numerical_features_all = [col for col in numerical_features_all if col
not in ['employee_id', 'salary']]

# Calculate correlation with target variable (salary)
correlations = df_engineered[numerical_features_all +
['salary']].corr()['salary'].abs().sort_values(ascending=False)
correlations = correlations.drop('salary') # Remove self-correlation

print("Feature Correlation with Salary (absolute values):")
print(correlations.round(3))

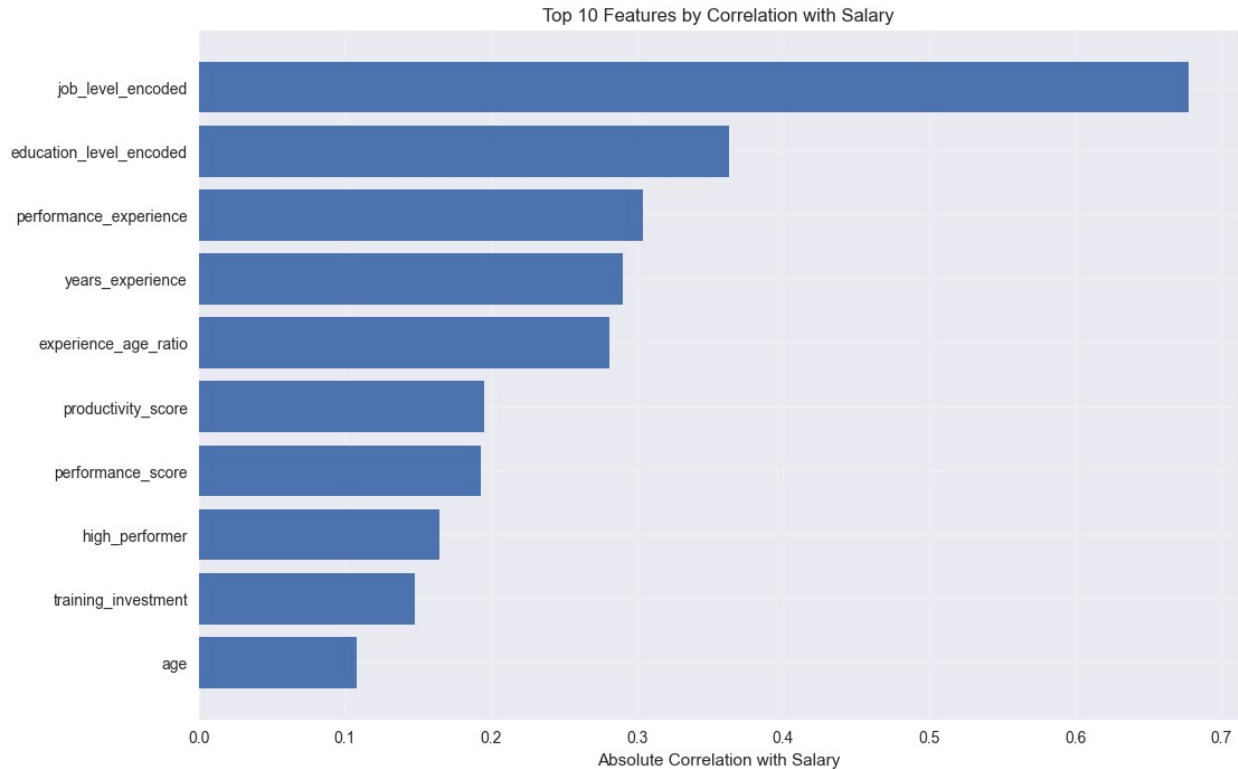
# Visualize feature importance
plt.figure(figsize=(12, 8))
top_features = correlations.head(10)
plt.barh(range(len(top_features)), top_features.values)
plt.yticks(range(len(top_features)), top_features.index)
plt.xlabel('Absolute Correlation with Salary')
plt.title('Top 10 Features by Correlation with Salary')
plt.gca().invert_yaxis()
plt.grid(True, alpha=0.3)
plt.show()

```

Feature Correlation with Salary (absolute values):

job_level_encoded	0.678
education_level_encoded	0.363
performance_experience	0.304
years_experience	0.290
experience_age_ratio	0.282
productivity_score	0.196
performance_score	0.193
high_performer	0.165
training_investment	0.148
age	0.108
training_hours	0.029
projects_completed	0.002
hours_per_week	0.002
work_intensity	0.002

Name: salary, dtype: float64



```
# Use Random Forest for feature importance
# Prepare features (only numerical for this example)
X = df_engineered[numerical_features_all]
y = df_engineered['salary']

# Train a Random Forest to get feature importance
rf = RandomForestRegressor(n_estimators=100, random_state=42)
rf.fit(X, y)

# Get feature importance
feature_importance = pd.DataFrame({
    'feature': X.columns,
    'importance': rf.feature_importances_
}).sort_values('importance', ascending=False)

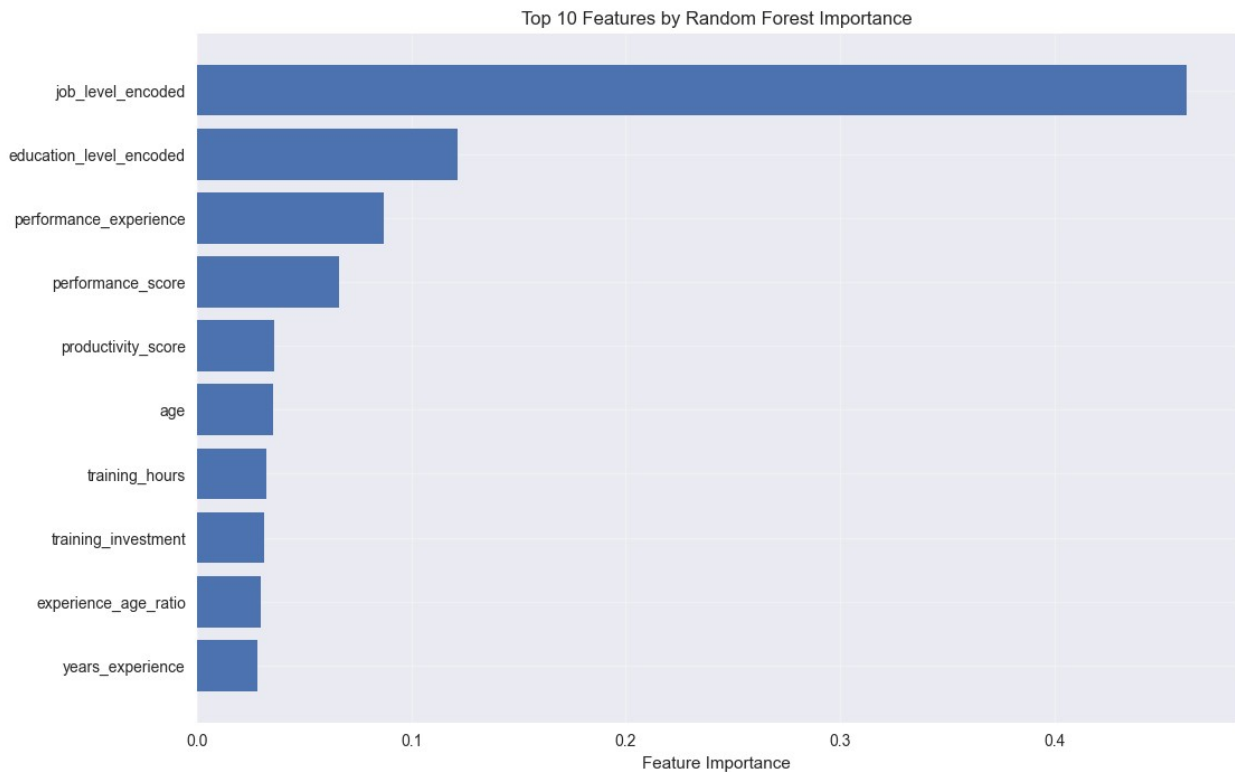
print("\nRandom Forest Feature Importance:")
print(feature_importance.round(4))

# Visualize feature importance
plt.figure(figsize=(12, 8))
top_rf_features = feature_importance.head(10)
plt.barh(range(len(top_rf_features)), top_rf_features['importance'])
plt.yticks(range(len(top_rf_features)), top_rf_features['feature'])
plt.xlabel('Feature Importance')
plt.title('Top 10 Features by Random Forest Importance')
plt.gca().invert_yaxis()
```

```
plt.grid(True, alpha=0.3)
plt.show()
```

Random Forest Feature Importance:

	feature	importance
7	job_level_encoded	0.4617
6	education_level_encoded	0.1216
12	performance_experience	0.0871
2	performance_score	0.0664
9	productivity_score	0.0363
0	age	0.0355
5	training_hours	0.0327
11	training_investment	0.0315
8	experience_age_ratio	0.0297
1	years_experience	0.0285
3	hours_per_week	0.0230
10	work_intensity	0.0222
4	projects_completed	0.0218
13	high_performer	0.0021



Step 8: Building a Preprocessing Pipeline

Let's create a reusable pipeline for all our preprocessing steps.

```

# Create a comprehensive preprocessing pipeline
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer

# Define column groups
numerical_cols = ['age', 'years_experience', 'performance_score',
                  'hours_per_week',
                  'projects_completed', 'training_hours']
categorical_cols = ['education_level', 'department', 'job_level',
                   'location']

# Create preprocessing pipelines for different column types
numerical_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])

categorical_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(drop='first', sparse_output=False))
])

# Combine pipelines
preprocessor = ColumnTransformer([
    ('num', numerical_pipeline, numerical_cols),
    ('cat', categorical_pipeline, categorical_cols)
])

print("Preprocessing Pipeline Created:")
print("\nNumerical Pipeline:")
print("  1. Impute missing values with median")
print("  2. Standard scaling (mean=0, std=1)")
print("\nCategorical Pipeline:")
print("  1. Impute missing values with mode")
print("  2. One-hot encoding (drop first category)")

Preprocessing Pipeline Created:

Numerical Pipeline:
  1. Impute missing values with median
  2. Standard scaling (mean=0, std=1)

Categorical Pipeline:
  1. Impute missing values with mode
  2. One-hot encoding (drop first category)

# Test the pipeline on our data
# Start with original data (with missing values)

```

```

X_original = df[numerical_cols + categorical_cols]
y_original = df['salary']

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X_original,
y_original, test_size=0.2, random_state=42)

print(f"Training set shape: {X_train.shape}")
print(f"Test set shape: {X_test.shape}")
print(f"\nMissing values in training set:
{X_train.isnull().sum().sum()}")
print(f"Missing values in test set: {X_test.isnull().sum().sum()}")

# Apply preprocessing pipeline
X_train_processed = preprocessor.fit_transform(X_train)
X_test_processed = preprocessor.transform(X_test)

print(f"\nAfter preprocessing:")
print(f"Training set shape: {X_train_processed.shape}")
print(f"Test set shape: {X_test_processed.shape}")
print(f"Missing values in processed training set:
{np.isnan(X_train_processed).sum()}")
print(f"Missing values in processed test set:
{np.isnan(X_test_processed).sum()}")

Training set shape: (802, 10)
Test set shape: (201, 10)

Missing values in training set: 113
Missing values in test set: 37

After preprocessing:
Training set shape: (802, 20)
Test set shape: (201, 20)
Missing values in processed training set: 0
Missing values in processed test set: 0

# Get feature names after preprocessing
# Numerical features keep their names
num_feature_names = numerical_cols

# Categorical features get expanded
cat_feature_names = preprocessor.named_transformers_['cat']
['onehot'].get_feature_names_out(categorical_cols)

all_feature_names = num_feature_names + list(cat_feature_names)

print(f"Total features after preprocessing: {len(all_feature_names)}")
print("\nFeature names:")
for i, name in enumerate(all_feature_names):
    print(f" {i+1:2d}. {name}")

```

Total features after preprocessing: 20

Feature names:

1. age
2. years_experience
3. performance_score
4. hours_per_week
5. projects_completed
6. training_hours
7. education_level_High School
8. education_level_Master
9. education_level_PhD
10. department_Finance
11. department_HR
12. department_Marketing
13. department_Sales
14. job_level_Lead
15. job_level_Mid
16. job_level_Senior
17. location_Chicago
18. location_New York
19. location_Remote
20. location_San Francisco

Challenge: Your Turn to Practice!

Now it's your turn to apply preprocessing techniques.

Challenge 1: Create a new feature

Create a feature that represents "career progression speed" (job level encoded divided by years of experience). Handle the case where years of experience is 0.

```
# Challenge 1: Create career progression speed feature
# Hint: Use the job_level_encoded and years_experience columns

# Create career progression speed feature
# Handle the case where years_experience is 0
df_engineered['career_progression_speed'] = np.where(
    df_engineered['years_experience'] == 0,
    0,
    df_engineered['job_level_encoded'] /
df_engineered['years_experience']
)

print("Career progression speed feature created")
print(df_engineered['career_progression_speed'].describe())
```

```
Career progression speed feature created
count      1003.000000
mean        1.149563
std         5.501231
min         0.000000
25%         0.000000
50%         0.209361
75%         0.664623
max         124.090519
Name: career_progression_speed, dtype: float64
```

Challenge 2: Handle a new categorical variable

Imagine we have a new column 'work_style' with values ['Remote', 'Hybrid', 'Office']. Add appropriate preprocessing for this column to our pipeline.

```
# Challenge 2: Add a new categorical variable to the preprocessing pipeline
# Hint: You'll need to modify the categorical_cols list and recreate the preprocessor

# Create the new 'work_style' column
np.random.seed(42)
df_with_work_style = df.copy()
df_with_work_style['work_style'] = np.random.choice(['Remote', 'Hybrid', 'Office'],
size=len(df_with_work_style))

# Add work_style to categorical columns
numerical_cols_updated = ['age', 'years_experience', 'performance_score', 'hours_per_week',
                           'projects_completed', 'training_hours']
categorical_cols_updated = ['education_level', 'department', 'job_level', 'location', 'work_style']

# Recreate the preprocessing pipeline
numerical_pipeline_updated = Pipeline([
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])

categorical_pipeline_updated = Pipeline([
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(drop='first', sparse_output=False))
])

preprocessor_updated = ColumnTransformer([
    ('num', numerical_pipeline_updated, numerical_cols_updated),
```



```

    ('cat', categorical_pipeline_updated, categorical_cols_updated)
])

print("Updated preprocessing pipeline with work_style column")
print(f"Categorical columns: {categorical_cols_updated}")

Updated preprocessing pipeline with work_style column
Categorical columns: ['education_level', 'department', 'job_level',
'location', 'work_style']

```

Challenge 3: Compare imputation strategies

Compare the effect of using mean vs median imputation for the 'performance_score' column. Which one preserves the distribution better?

```

# Challenge 3: Compare mean vs median imputation for performance_score
# Hint: Create two versions of the data with different imputation
# strategies and compare histograms

# Create two versions with different imputation
df_mean_imputed = df.copy()
df_median_imputed = df.copy()

# Mean imputation
mean_value = df_mean_imputed['performance_score'].mean()
df_mean_imputed['performance_score'].fillna(mean_value, inplace=True)

# Median imputation
median_value = df_median_imputed['performance_score'].median()
df_median_imputed['performance_score'].fillna(median_value,
inplace=True)

# Compare distributions
fig, axes = plt.subplots(1, 3, figsize=(15, 4))

axes[0].hist(df['performance_score'].dropna(), bins=20, alpha=0.7)
axes[0].set_title('Original Distribution')
axes[0].set_xlabel('Performance Score')


axes[1].hist(df_mean_imputed['performance_score'], bins=20, alpha=0.7)
axes[1].set_title('After Mean Imputation')
axes[1].set_xlabel('Performance Score')

axes[2].hist(df_median_imputed['performance_score'], bins=20,
alpha=0.7)
axes[2].set_title('After Median Imputation')
axes[2].set_xlabel('Performance Score')

plt.tight_layout()
plt.show()

```

```
print(f"Mean imputation value: {mean_value:.3f}")
print(f"Median imputation value: {median_value:.3f}")
print("\nMedian imputation better preserves the distribution as it is less affected by outliers")
```



```
Mean imputation value: 7.371
Median imputation value: 7.395
```

Median imputation better preserves the distribution as it is less affected by outliers

Summary

Congratulations! You've mastered the essential data preprocessing and feature engineering techniques. Here's what you've learned:

☐ Key Skills Mastered:

1. **Data Quality Assessment:** Identifying and understanding missing values, duplicates, and outliers
2. **Missing Value Handling:** Multiple imputation strategies (mean, median, mode, group-based)
3. **Categorical Encoding:** Ordinal encoding for ordered categories, one-hot encoding for nominal categories
4. **Feature Scaling:** StandardScaler and MinMaxScaler for different use cases
5. **Outlier Treatment:** Detection using IQR method and treatment using Winsorization
6. **Feature Engineering:** Creating new features from existing ones using domain knowledge
7. **Feature Selection:** Using correlation and Random Forest importance for feature ranking
8. **Pipeline Creation:** Building reusable preprocessing pipelines with scikit-learn

☐ Best Practices Learned:

- Always understand your data before preprocessing
- Choose imputation strategies based on data distribution and missing patterns
- Use appropriate encoding for different types of categorical variables
- Scale features when algorithms are sensitive to feature magnitude

- Handle outliers based on domain knowledge and model requirements
- Create features that capture domain expertise
- Use pipelines to ensure reproducibility and prevent data leakage

□ Next Steps:

In the next lab, we'll use this preprocessed data to build and evaluate classification models, learning how to:

- Train logistic regression and decision tree classifiers
- Evaluate model performance using various metrics
- Use cross-validation for robust model assessment
- Interpret model results and make predictions

□ Additional Resources:

- [Scikit-learn Preprocessing Guide](#)
- [Feature Engineering Techniques](#)
- [Handling Missing Data](#)
- [Pipeline Documentation](#)